

## Spark2.3 原生支持用 Kubernetes 调度

The diagram illustrates the Spark-Kubernetes integration architecture, showing the flow of communication between Spark and Kubernetes components.

**Spark Side (Left):**

- The process starts with the `spark-submit` command (1).
- This command runs within the **Spark** environment (2).
- Inside Spark, the `KubernetesClientApplication` is executed (3).
- The `KubernetesClientApplication` interacts with the `Client` (4).
- The `Client` interacts with the `KubernetesClient` (5).

**Kubernetes Side (Right):**

- The `KubernetesClient` from Spark connects to the `api-server` (6).
- The `api-server` interacts with the `KubernetesDriverEndpoint` (7).
- The `api-server` also interacts with the `KubernetesClientSchedulerBackend` (8).
- The `api-server` interacts with the `KubernetesClient` (9).
- The `api-server` interacts with the `CoarseGrainedExecutorBackend` (10).
- The `api-server` interacts with the `CoarseGrainedExecutorBackend` (11).

**Internal Components:**

- Pod:** The outermost container.
- Container:** The container within the pod.
- Driver:** The driver component within the container.
- SparkContext:** The SparkContext component within the driver.
- KubernetesClusterManager:** The KubernetesClusterManager component within the SparkContext.
- KubernetesClientSchedulerBackend:** The KubernetesClientSchedulerBackend component within the KubernetesClusterManager.
- KubernetesDriverEndpoint:** The KubernetesDriverEndpoint component within the KubernetesClientSchedulerBackend.
- KubernetesClient:** The KubernetesClient component within the KubernetesClientSchedulerBackend.
- Pod:** The outermost container.
- Container:** The container within the pod.
- Container:** The container within the pod.
- CoarseGrainedExecutorBackend:** The CoarseGrainedExecutorBackend component within the container.

启动 `KubernetesClientApplication` 之后，创建一个包含能够与 Kubernetes 的 `api-server` 通信的 `KubernetesClient` 的 `Client`，随后，在根据 `SparkConf` 拼接完成用于在 Kubernetes 中启动 `Driver` 的 `yaml` 文件之后，`Client` 会通知 `api-server` 在 Kubernetes 集群中启动一个包含一个用于启动 `Driver` 的 `Container` 的 `Pod`。

初始化完成后，TaskScheduler 会启动，同时会启动 KubernetesClientSchedulerBackend。当 KubernetesClientSchedulerBackend 启动的时候，创建一个 KubernetesDriverEndpoint 和一个用于和 api-server 进行通讯的 KubernetesClient。

KubernetesClientSchedulerBackend 如果开启了动态分配，那么会定期去检查当前所需要的 executor 的个数，如果没有开启动态分配，那么会将所需的 executor 个数设置为 SparkConf 中所设定的数目。当检查到正在运行的包含 executor 少于所需要的 executor 的数目，那么会启动缺少的 executor 的数目个 Pod，每个 Pod 包含一个 executor，如果缺少的 executor 数目比一次能够启动的 executor 数目还要多，那么只启动一次最多能启动的数目个 executor。KubernetesClientSchedulerBackend 利用之前所创建的 KubernetesClient 通知 api-server 去启动 Pod。Pod 在启动的时候，利用 KubernetesClientSchedulerBackend 根据 SparkConf 所拼接的 yaml 文件，来创建 Pod，然后启动 CoarseGrainedExecutorBackend 以创建、管理、销毁 executor，并使 executor 与 driver 互相通信。

## 源码分析

### spark-submit

```
if [ -z "${SPARK_HOME}" ]; then
    source "$(dirname "$0")/find-spark-home
fi

# disable randomized hash for string in Python 3.3+
export PYTHONHASHSEED=0

exec "${SPARK_HOME}"/bin/spark-class org.apache.spark.deploy.SparkSubmit "$@"
```

spark-submit 通过 spark-class 来运行 org.apache.spark.deploy.SparkSubmit 这个程序，其参数传递给 SparkSubmit。

### SparkSubmit

```
override def main(args: Array[String]): Unit = {
    // ...
    val appArgs = new SparkSubmitArguments(args)
    // ...
    appArgs.action match {
        case SparkSubmitAction.SUBMIT => submit(appArgs, uninitLog)
        case SparkSubmitAction.KILL => kill(appArgs)
        case SparkSubmitAction.REQUEST_STATUS => requestStatus(appArgs)
    }
}

private def submit(args: SparkSubmitArguments, uninitLog: Boolean): Unit = {
    val (childArgs, childClasspath, sparkConf, childMainClass) = prepareSubmitEnvironment(args)

    def doRunMain(): Unit = {
        if (args.proxyUser != null) {
            val proxyUser = UserGroupInformation.createProxyUser(args.proxyUser,
                UserGroupInformation.getCurrentUser())
            try {
                proxyUser.doAs(new PrivilegedExceptionAction[Unit]() {
                    override def run(): Unit = {
                        runMain(childArgs, childClasspath, sparkConf, childMainClass, args.verbose)
                    }
                })
            }
        }
    }
}
```

```

    })
  } catch {...}
} else {
  runMain(childArgs, childClasspath, sparkConf, childMainClass, args.verbose)
}
}
// ...
if (args.isStandaloneCluster && args.useRest) {
  try {
    printStream.println("Running Spark using the REST application submission protocol.")
    doRunMain()
  } catch {
    case e: SubmitRestConnectionException => //...
      args.useRest = false
      submit(args, false)
  }
} else {
  doRunMain()
}
}

private def runMain(
  childArgs: Seq[String],
  childClasspath: Seq[String],
  sparkConf: SparkConf,
  childMainClass: String,
  verbose: Boolean): Unit = {
  // ...
  var mainClass: Class[_] = null

  try {
    mainClass = Utils.classForName(childMainClass)
  } catch {...}

  val app: SparkApplication = if (classOf[SparkApplication].isAssignableFrom(mainClass)) {
    mainClass.newInstance().asInstanceOf[SparkApplication]
  } else {
    // SPARK-4170
    if (classOf[scala.App].isAssignableFrom(mainClass)) {
      printWarning("Subclasses of scala.App may not work correctly. Use a main() method
instead.")
    }
    new JavaMainApplication(mainClass)
  }
  // ...
  try {
    app.start(childArgs.toArray, sparkConf)
  } catch {...}
}

```

SparkSubmit 首先解析传入进来的参数，从中得到 childArgs, childClasspath, sparkConf, childMainClass 四个主要的参数信息，然后启动 childMainClass，在 Spark on Kubernetes 中，childMainClass 就是 org.apache.spark.deploy.k8s.submit.KubernetesClientApplication。因此，SparkSubmit 启动 KubernetesClientApplication。

## KubernetesClientApplication

```
private def run(clientArguments: ClientArguments, sparkConf: SparkConf): Unit = {
  val namespace = sparkConf.get(KUBERNETES_NAMESPACE)
  val kubernetesAppId = s"spark-${UUID.randomUUID().toString.replaceAll("-", "")}"
  val launchTime = System.currentTimeMillis()
  val waitForAppCompletion = sparkConf.get(WAIT_FOR_APP_COMPLETION)
  val appName = sparkConf.getOption("spark.app.name").getOrElse("spark")
  val master = sparkConf.get("spark.master").substring("k8s://".length)
  val loggingInterval = if (waitForAppCompletion) Some(sparkConf.get(REPORT_INTERVAL)) else None

  val watcher = new LoggingPodStatusWatcherImpl(kubernetesAppId, loggingInterval)

  val orchestrator = new DriverConfigOrchestrator(
    kubernetesAppId,
    launchTime,
    clientArguments.mainAppResource,
    appName,
    clientArguments.mainClass,
    clientArguments.driverArgs,
    sparkConf)

  Utils.tryWithResource(SparkKubernetesClientFactory.createKubernetesClient(
    master,
    Some(namespace),
    KUBERNETES_AUTH_SUBMISSION_CONF_PREFIX,
    sparkConf,
    None,
    None)) { kubernetesClient =>
    val client = new Client(
      orchestrator.getAllConfigurationSteps,
      sparkConf,
      kubernetesClient,
      waitForAppCompletion,
      appName,
      watcher)
    client.run()
  }
}
```

KubernetesClientApplication 启动时，首先再次解析传入的参数，之后调用上面的 run 方法来启动，然后创建一个包含一个可以与 Kubernetes 的 api-server 通信的 KubernetesClient 的 Client。之后，启动这个 client。

```
def run(): Unit = {
  var currentDriverSpec = KubernetesDriverSpec.initialSpec(sparkConf)
  for (nextStep <- submissionSteps) {
    currentDriverSpec = nextStep.configureDriver(currentDriverSpec)
  }
  val resolvedDriverJavaOpts = currentDriverSpec
    .driverSparkConf
    .remove(org.apache.spark.internal.config.DRIVER_JAVA_OPTIONS)
```

```

        .getAll
        .map {
            case (confKey, confValue) => s"-D$confKey=$confValue"
        } ++ driverJavaOptions.map(Utils.splitCommandString).getOrElse(Seq.empty)
    val driverJavaOptsEnvs: Seq[EnvVar] = resolvedDriverJavaOpts.zipWithIndex.map {
        case (option, index) =>
            new EnvVarBuilder()
                .withName(s"$ENV_JAVA_OPT_PREFIX$index")
                .withValue(option)
                .build()
    }
    val resolvedDriverContainer = new ContainerBuilder(currentDriverSpec.driverContainer)
        .addAllToEnv(driverJavaOptsEnvs.asJava)
        .build()
    val resolvedDriverPod = new PodBuilder(currentDriverSpec.driverPod)
        .editSpec()
        .addToContainers(resolvedDriverContainer)
        .endSpec()
        .build()
    Utils.tryWithResource(
        kubernetesClient
            .pods()
            .withName(resolvedDriverPod.getMetadata.getName)
            .watch(watcher)) { _ =>
        val createdDriverPod = kubernetesClient.pods().create(resolvedDriverPod)
        try {
            if (currentDriverSpec.otherKubernetesResources.nonEmpty) {
                val otherKubernetesResources = currentDriverSpec.otherKubernetesResources
                addDriverOwnerReference(createdDriverPod, otherKubernetesResources)
                kubernetesClient.resourceList(otherKubernetesResources: _*).createOrReplace()
            }
        } catch {
            case NonFatal(e) =>
                kubernetesClient.pods().delete(createdDriverPod)
                throw e
        }
        if (waitForAppCompletion) {
            logInfo(s"Waiting for application $appName to finish...")
            watcher.awaitCompletion()
            logInfo(s"Application $appName finished.")
        } else {
            logInfo(s"Deployed Spark application $appName into Kubernetes.")
        }
    }
}

```

Client 启动之后，在拼接了用于配置启动 Driver 的 yaml 文件之后，利用之前创建的 KubernetesClient，通知 api-server 启动 Pod，这个 Pod 中包含一个 container，在这个 container 中，启动 Driver。其中拼接出来的 yaml 文件中包含在 spark-submit 时所提交的参数，这些参数中包含了 Pod 中所应该启动的程序是什么程序，每个 Pod 需要哪些资源等。

在 Pod 中启动了用户所提交的应用程序之后，启动 SparkContext。

## SparkContext

```
try {
  //...
  val (sched, ts) = SparkContext.createTaskScheduler(this, master, deployMode)
  _schedulerBackend = sched
  _taskScheduler = ts
  _dagScheduler = new DAGScheduler(this)
  _heartbeatReceiver.ask[Boolean](TaskSchedulerIsSet)

  // start TaskScheduler after taskScheduler sets DAGScheduler reference in DAGScheduler's
  // constructor
  _taskScheduler.start()
  //...
  // Optionally scale number of executors dynamically based on workload. Exposed for testing.
  val dynamicAllocationEnabled = Utils.isDynamicAllocationEnabled(_conf)
  _executorAllocationManager =
    if (dynamicAllocationEnabled) {
      schedulerBackend match {
        case b: ExecutorAllocationClient =>
          Some(new ExecutorAllocationManager(
            schedulerBackend.asInstanceOf[ExecutorAllocationClient], listenerBus, _conf,
            _env.blockManager.master))
        case _ =>
          None
      }
    } else {
      None
    }
  _executorAllocationManager.foreach(_.start())
  //...
  setupAndStartListenerBus()
  postEnvironmentUpdate()
  postApplicationStart()
  //...
} catch {...}

private def createTaskScheduler(
  sc: SparkContext,
  master: String,
  deployMode: String): (SchedulerBackend, TaskScheduler) = {
  import SparkMasterRegex._
  // When running locally, don't try to re-execute tasks on failure.
  val MAX_LOCAL_TASK_FAILURES = 1
  master match {
    case "local" => //...
    case LOCAL_N_REGEX(threads) => //...
    case LOCAL_N_FAILURES_REGEX(threads, maxFailures) => //...
    case SPARK_REGEX(sparkUrl) => //...
    case LOCAL_CLUSTER_REGEX(numSlaves, coresPerSlave, memoryPerSlave) => //...
    case masterUrl =>

      val cm = getClusterManager(masterUrl) match {
```

```

        case Some(clusterMgr) => clusterMgr
        case None => throw new SparkException("Could not parse Master URL: '" + master + "'")
    }
    try {
        val scheduler = cm.createTaskScheduler(sc, masterUrl)
        val backend = cm.createSchedulerBackend(sc, masterUrl, scheduler)
        cm.initialize(scheduler, backend)
        (backend, scheduler)
    } catch {...}
}
}

```

在 SparkContext 中，创建 SparkContext 的同时就开始创建 TaskScheduler。首先获取对应的模式的集群资源管理器，在 Spark on kubernetes 中获取的是 KubernetesClusterManager。之后 KubernetesClusterManager 创建一个 TaskScheduler，并利用这个 TaskScheduler 创建 KubernetesClientSchedulerBackend。之后初始化 TaskScheduler 和 KubernetesClientSchedulerBackend。然后 SparkContext 会利用 TaskScheduler 创建一个 DAGScheduler。DAGScheduler 负责把用户提交的计算任务划分为一个个调度阶段。然后启动 TaskScheduler。

TaskScheduler 的启动过程，首先会把 KubernetesClientSchedulerBackend 同时也启动，然后 KubernetesClientSchedulerBackend 会启动 KubernetesDriverEndpoint 和 另外一个KubernetesClient。以此使得 Driver 可以控制 Kubernetes 集群和 Spark 集群，一方面通过 KubernetesClient 可以获取与 Kubernetes 集群沟通的能力，可以在需要的时候创建或者删除 Pod，另一方面通过 KubernetesDriverEndpoint 与 Spark 集群内的其它 executor 进行沟通，获取 executor 计算结果，删除无用的 executor 等。

## KubernetesClientSchedulerBackend

```

override def start(): Unit = {
    super.start()
    executorWatchResource.set(
        kubernetesClient
        .pods()
        .withLabel(SPARK_APP_ID_LABEL, applicationId())
        .watch(new ExecutorPodsWatcher()))

    allocatorExecutor.scheduleWithFixedDelay(
        allocatorRunnable, 0L, podAllocationInterval, TimeUnit.MILLISECONDS)
    if (!Utils.isDynamicAllocationEnabled(conf)) {
        doRequestTotalExecutors(initialExecutors)
    }
}

private val allocatorRunnable = new Runnable {
    private val executorReasonCheckAttemptCounts = new mutable.HashMap[String, Int]
    override def run(): Unit = {
        handleDisconnectedExecutors()
        val executorsToAllocate = mutable.Map[String, Pod]()
        val currentTotalRegisteredExecutors = totalRegisteredExecutors.get
        val currentTotalExpectedExecutors = totalExpectedExecutors.get
        val currentNodeToLocalTaskCount = getNodesWithLocalTaskCounts()
        RUNNING_EXECUTOR_PODS_LOCK.synchronized {
            if (currentTotalRegisteredExecutors < runningExecutorsToPods.size) {
                logDebug("Waiting for pending executors before scaling")
            }
        }
    }
}

```

```

    } else if (currentTotalExpectedExecutors <= runningExecutorsToPods.size) {
        logDebug("Maximum allowed executor limit reached. Not scaling up further.")
    } else {
        for (_ <- 0 until math.min(
            currentTotalExpectedExecutors - runningExecutorsToPods.size, podAllocationSize)) {
            val executorId = EXECUTOR_ID_COUNTER.incrementAndGet().toString
            val executorPod = executorPodFactory.createExecutorPod(
                executorId,
                applicationId(),
                driverUrl,
                conf.getExecutorEnv,
                driverPod,
                currentNodeToLocalTaskCount)
            executorsToAllocate(executorId) = executorPod
            logInfo(
                s"Requesting a new executor, total executors is now ${runningExecutorsToPods.size}")
        }
    }
}

val allocatedExecutors = executorsToAllocate.mapValues { pod =>
    Utils.tryLog {
        kubernetesClient.pods().create(pod)
    }
}

RUNNING_EXECUTOR_PODS_LOCK.synchronized {
    allocatedExecutors.map {
        case (executorId, attemptedAllocatedExecutor) =>
            attemptedAllocatedExecutor.map { successfullyAllocatedExecutor =>
                runningExecutorsToPods.put(executorId, successfullyAllocatedExecutor)
            }
    }
}

// ...
}

```

在启动 `KubernetesClientSchedulerBackend` 时，首先启动 `KubernetesDriverEndpoint`，之后设置周期性调度 executor，如果启动了动态调度，则每次获取当前所需的 executor 个数，否则，所需的 executor 个数是一个在 `SparkConf` 中确定的数目。与当前正在运行的 executor 数目进行对比，如果少了，则启动少了的数目个 Pod，这些 Pod 中包含一个 executor，或启动一次最大能启动的 Pod 数目个 Pod。

## entrypoint.sh

```

case "$SPARK_K8S_CMD" in
    driver)
        CMD=(
            ${JAVA_HOME}/bin/java
            "${SPARK_JAVA_OPTS[@]}"
            -cp "$SPARK_CLASSPATH"
            -Xms$SPARK_DRIVER_MEMORY
            -Xmx$SPARK_DRIVER_MEMORY
            -Dspark.driver.bindAddress=$SPARK_DRIVER_BIND_ADDRESS

```



```

    $SPARK_DRIVER_CLASS
    $SPARK_DRIVER_ARGS
)
;;

executor)
    CMD=(
        ${JAVA_HOME}/bin/java
        "${SPARK_JAVA_OPTS[@]}"
        -Xms$SPARK_EXECUTOR_MEMORY
        -Xmx$SPARK_EXECUTOR_MEMORY
        -cp "$SPARK_CLASSPATH"
        org.apache.spark.executor.CoarseGrainedExecutorBackend
        --driver-url $SPARK_DRIVER_URL
        --executor-id $SPARK_EXECUTOR_ID
        --cores $SPARK_EXECUTOR_CORES
        --app-id $SPARK_APPLICATION_ID
        --hostname $SPARK_EXECUTOR_POD_IP
    )
    ;;

init)
    CMD=(
        "$SPARK_HOME/bin/spark-class"
        "org.apache.spark.deploy.k8s.SparkPodInitContainer"
        "$@"
    )
    ;;

*)
    echo "Unknown command: $SPARK_K8S_CMD" 1>&2
    exit 1
esac

```

当 Kubernetes 集群的 api-server 收到创建 Pod 的消息时，会根据 yaml 文件启动对应 docker 的镜像，并且根据 entrypoint.sh 来启动程序。Driver 对应的是用户自己所提交的程序，Executor 所对应的则是 org.apache.spark.executor.CoarseGrainedExecutorBackend。在这个程序中，会启动一个 RpcEndpoint 以此来与 Driver 的 KubernetesClientEndpoint 通信。